



Traceless Execution Support for Privacy Enhancing Technologies

Daniela Lopes

Advisor: Prof. Nuno Santos

24 November 2021



Presentation outline

1. Introduction and motivation
2. Related work
3. Design
4. Evaluation
5. Conclusions



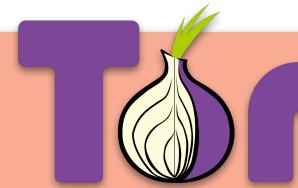
Why people need to use PETs

Privacy Enhancing Technologies (PETs):



Secure Messaging

VPNs



Anonymizing networks

Anti tracking



People use PETs to overcome several challenges:



Censorship



Lack of privacy



Surveillance



However, people are afraid to use PETs

Drawback of using PETs:

- Large amount of local resources
 - Non-trivial meta-data and cryptographic artifacts
- 
- Extensive digital footprint**

Why existing PETs are not enough to protect users:

- Forensic inspection of the devices of endangered users such as journalists, whistleblowers and travelers passing border controls
- PETs do not provide plausible deniability



Example of a threatening usage scenario

John passes border control to enter the repressive country and has his device inspected

→ **Snapshot 1**



John makes publications about his sensitive findings via the Tor browser on Wikileaks

John passes border control to leave the repressive country and has his device inspected a second time

→ **Snapshot 2**



The investigator compares **Snapshot 1** and **Snapshot 2** and observes that John was the author of the publications



Our goal: plausible deniable PET sessions

John passes border control to enter the repressive country and has his device inspected

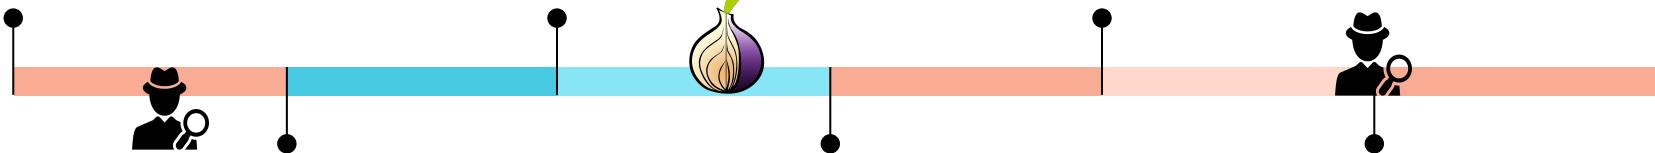
→ Snapshot 1

John makes publications about his sensitive findings via the Tor browser on WikiLeaks



John passes border control to leave the repressive country and has his device inspected a second time

→ Snapshot 2



John initiates a plausible deniable PET session

John closes the plausible deniable PET session

The investigator compares **Snapshot 1** and **Snapshot 2** and **cannot observe** changes that lead to suspecting John

Plausible deniability: ability to deny knowledge or responsibility of having installed and executed PETs on their device



Challenges of enforcing deniable PET sessions

Challenges:

1. Where is the PET data going to be stored?
2. How to make that data non observable?
3. How to maintain the PET functionality?



Contributions

Design: Calypso, a steganographic storage system

- Offer the abstraction of a hidden **shadow partition**, to be used as a regular partition
- The shadow partition parasites on the free blocks of the native system to conceal data
- Preserves the initial entropy characteristics of the disk

Implementation: Fully functional prototype as a Linux kernel module for version 5.4

Evaluation: Extensive evaluation of the prototype

- Functionality
- Isolation
- Security
- Performance



Related work

Private program execution

- ➔ Processes are isolated from the rest of the system, preventing leaks
- ➔ Not deniable by design, the system itself is not concealed

Deniable file systems

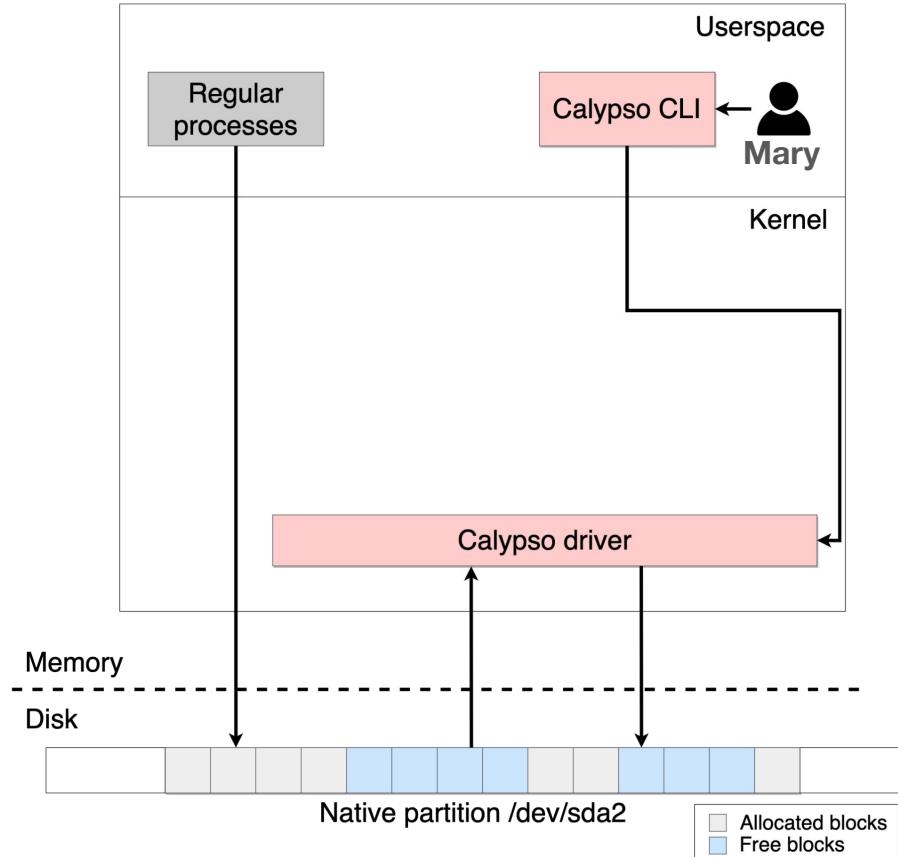
- ➔ Data hiding in persistent storage
- ➔ Writing random patterns abnormally increases the disk's entropy, disclosing the existence of hidden data

Defending from multi-snapshot attacks

- ➔ Focus on preventing multi-snapshot attack, thereby providing stronger deniability
- ➔ Performance penalties and unnatural disk write-patterns

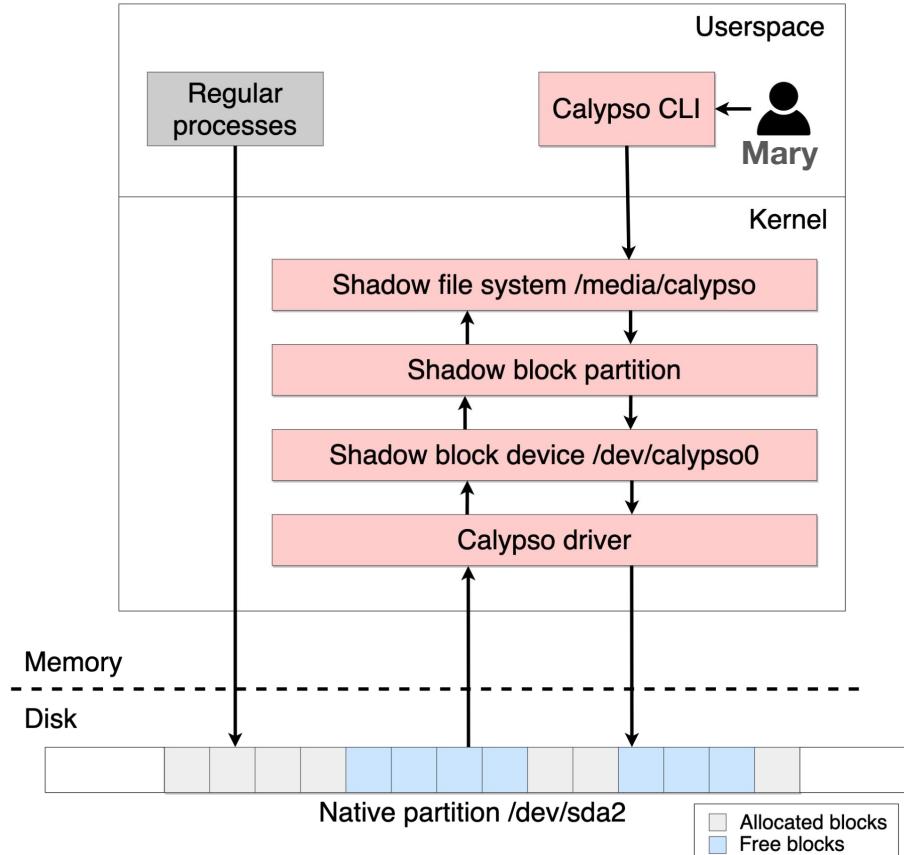


Calypso's architecture: inserting the module



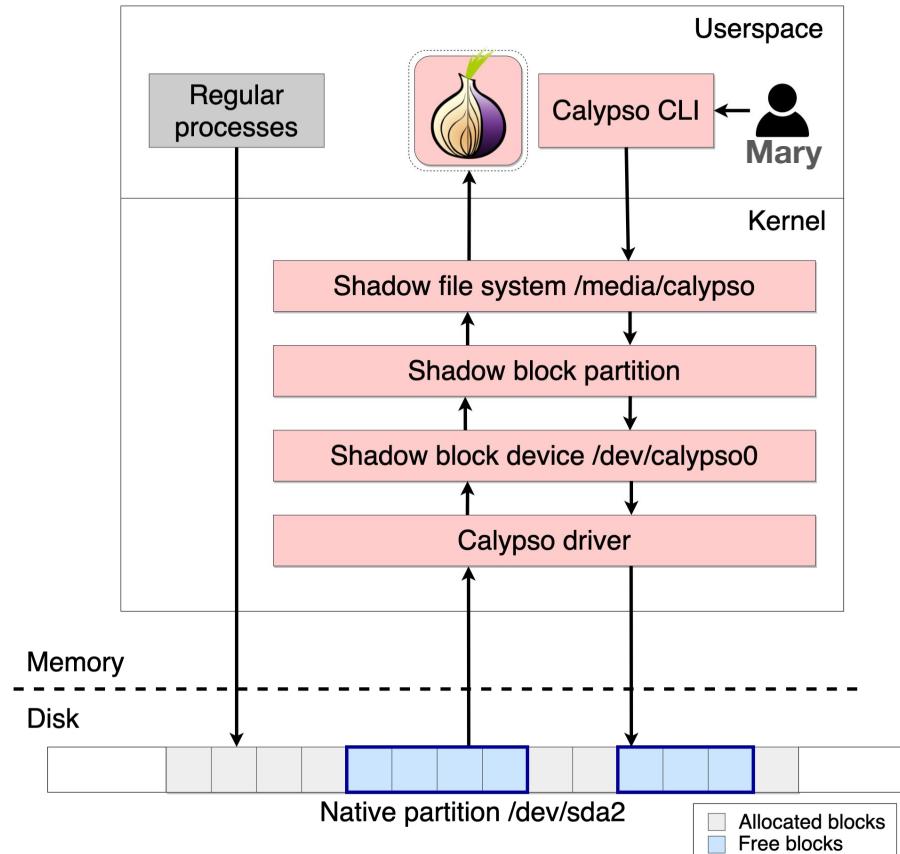


Calypso's architecture: shadow partition



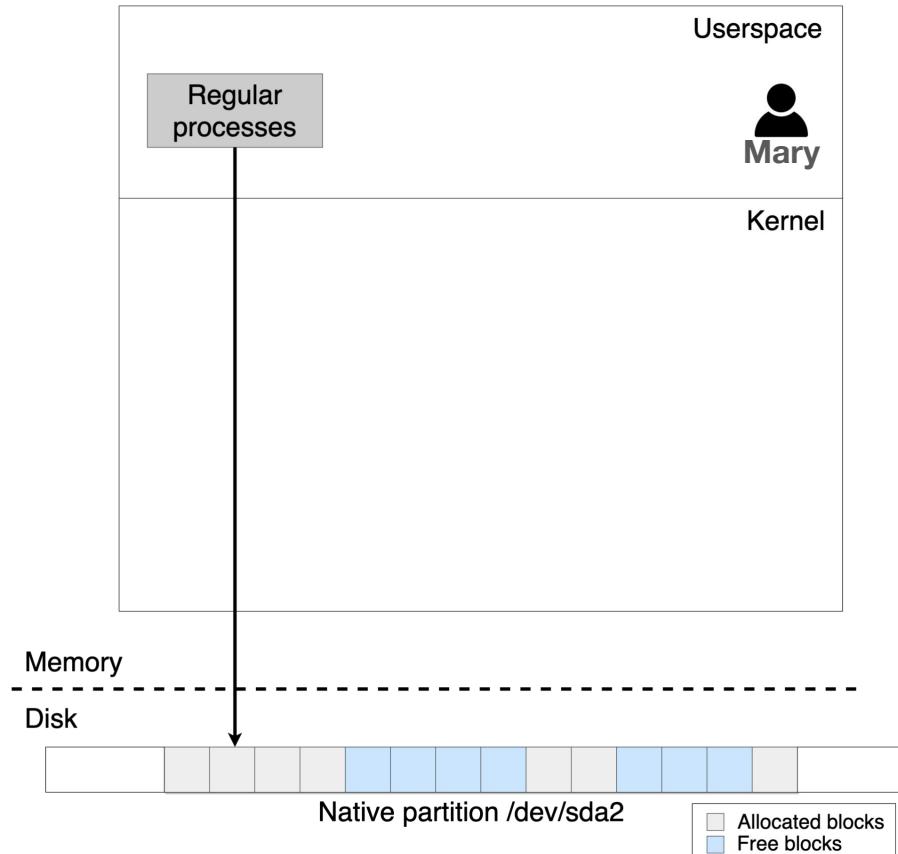


Calypso's architecture: executing PETs





Calypso's architecture: state after Calypso





Calypso: design obstacles

Traceless bootstrap

- Prevent or erase traces resulted from executing Calypso and inserting its module

Block allocation

- Find usable blocks without disrupting the native system or causing corruption to Calypso

Monitoring native changes to blocks

- Track free blocks and prevent overwrites to Calypso data

Deniably encoding data blocks

- Ensure Calypso only makes deniable changes to native blocks



Calypso: design solutions

Traceless bootstrap

- Erase all persistent traces of bootstrap using rootkit techniques

Block allocation

- Use the free blocks of the native file system, mapped with in-memory data structures

Monitoring native changes to blocks

- Intercept all write requests to the native system to check and update the allocation in-memory data structures accordingly

Deniably encoding data blocks

- Choose only blocks with initial entropy above an established threshold



Evaluation goals

Functionality: Can Calypso support the execution of multiple PETs without disrupting the system or disturbing the functionality of PETs?

Isolation: Does the native file system present detectable changes after executing Calypso?

Security: Does Calypso make only deniable changes that preserve the initial entropy of the disk while ensuring practical storage capacity?

Performance: Is Calypso performant enough to maintain the usability of PETs and the native system?

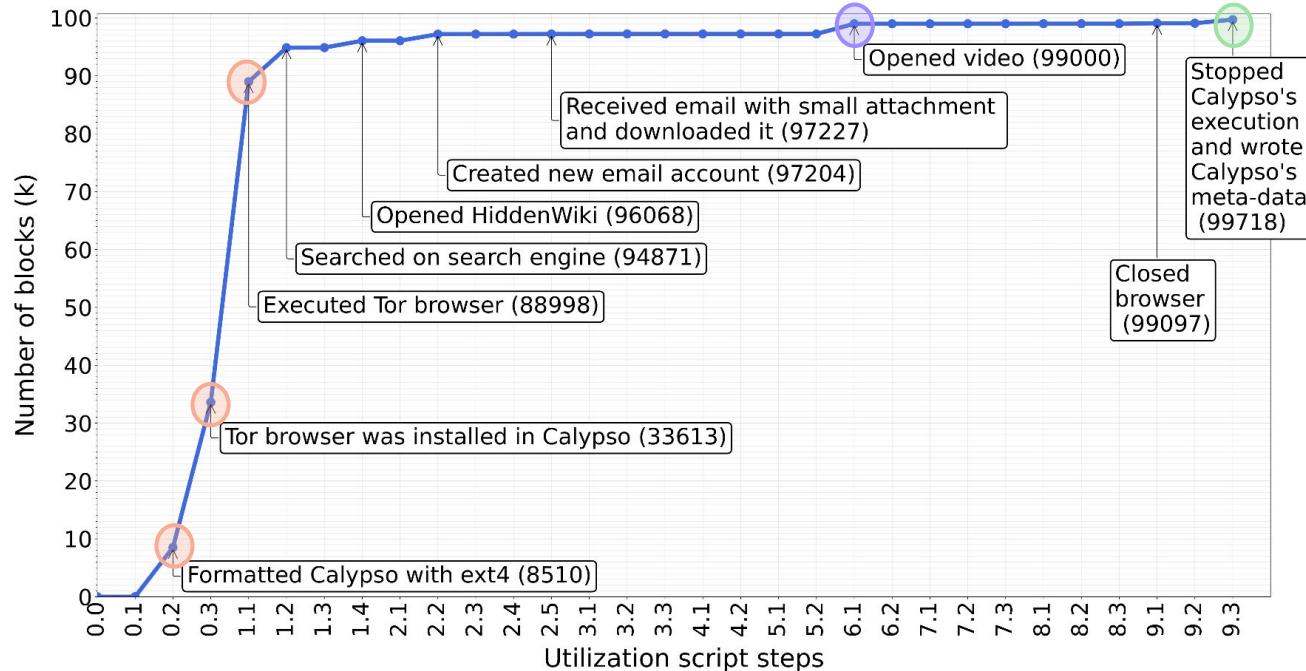


Functionality evaluation

- Test Calypso's ability to sustain the execution of PET applications without file system errors
- Monitor the execution of PETs with multiple disk workloads, throughout the execution of multiple utilization steps



Functionality results



Calypso is able to sustain the execution of PETs through multiple utilization steps, without disrupting the system or the PET

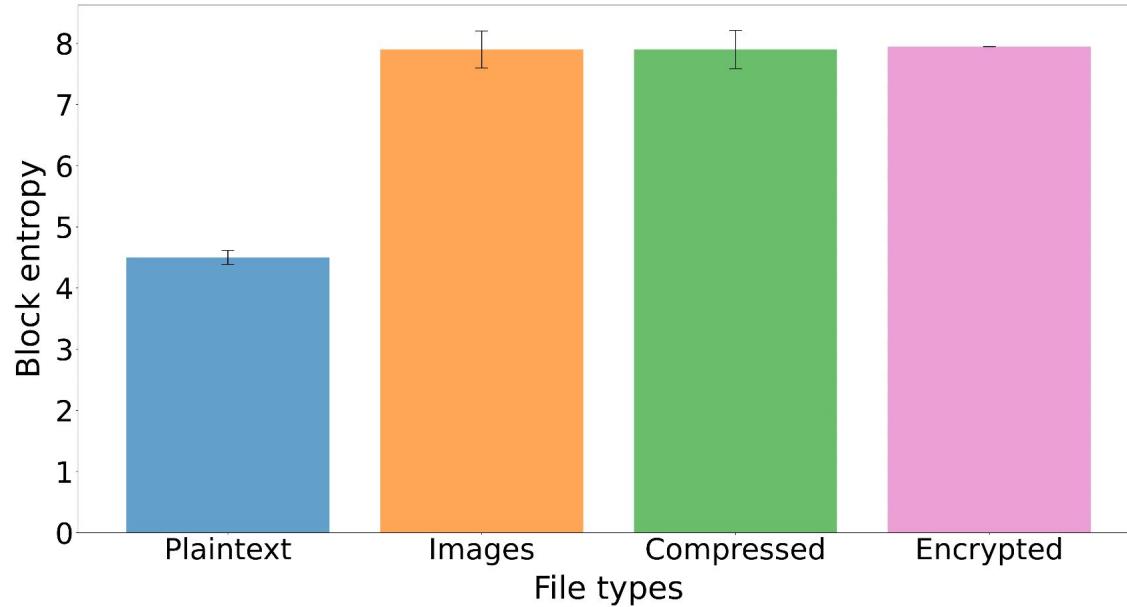


Security evaluation

- Assess Calypso's capability to make only non-observable changes to the native file system
- Maintain a practical storage capacity to execute PETs
- During a physical inspection, the inspector should not be able to distinguish between free blocks that store Calypso's data and meta-data, from blocks that do not



Security results: entropy of common files



Several files erased in our computers can result in usable blocks for Calypso, without changing the initial entropy characteristics



Security

Entropy threshold: all blocks with entropy above this value are eligible to be used

Entropy difference: quantifies how much the entropy of a block has changed

Entropy threshold	Usable blocks (%)	Storage capacity (GB)	Maximum entropy difference
0	100%	4.00	1.000
5	74.91%	3.00	0.372
7	74.67%	2.99	0.120

Calypso can produce only deniable changes with higher entropy values while preserving sufficient storage capacity



Conclusions

- State-of-the-art PETs are vulnerable to forensic analysis
- Calypso makes non-observable changes to disk
- Give protection against forensic analysis

Future work:

- Enhance isolation by leveraging containerization techniques
- Explore the implications of purposely executing plausible actions as a decoy to justify the changes to disk and to increase storage capacity



Q&A



Extra slides



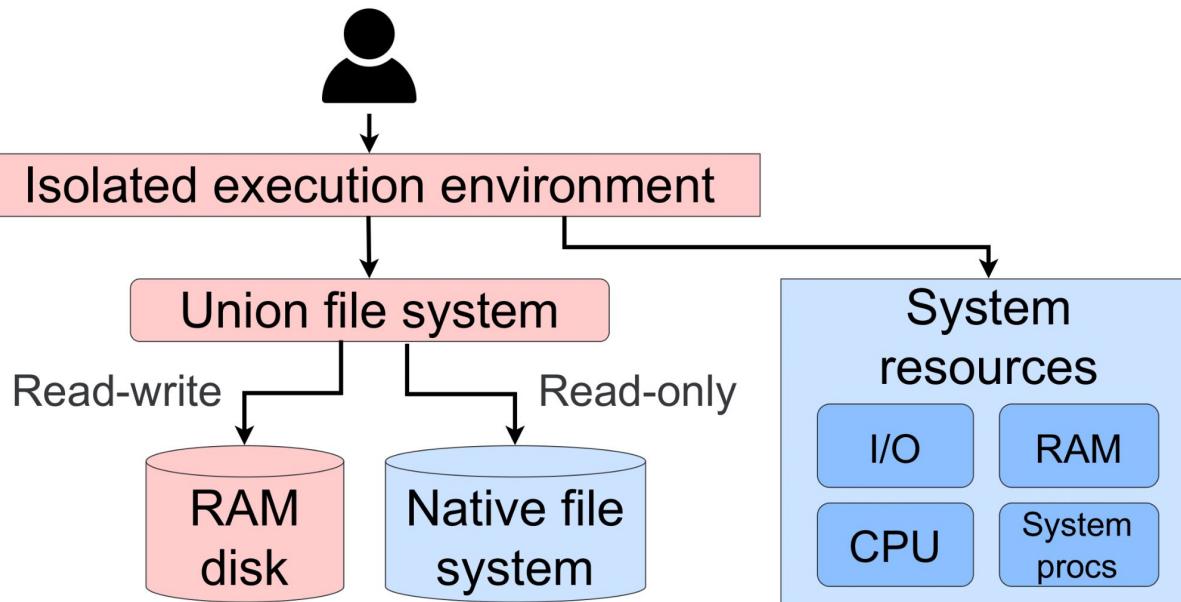
Main findings

Functionality: Calypso can sustain the execution of PETs without disrupting the native system or being corrupted

Security: Calypso can take advantage of several native blocks generated by the erasure of common file types to make only deniable changes, while preserving sufficient storage capacity



Residue-free computing



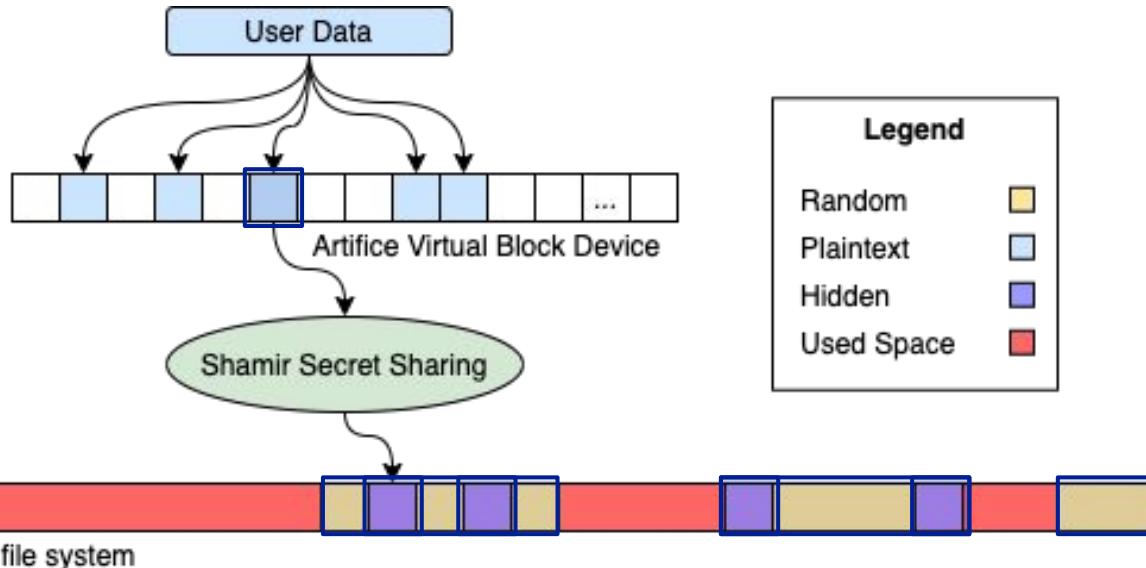
No persistence
or plausible
deniability

Lack of
application
support

Changes the file
system



Artifice

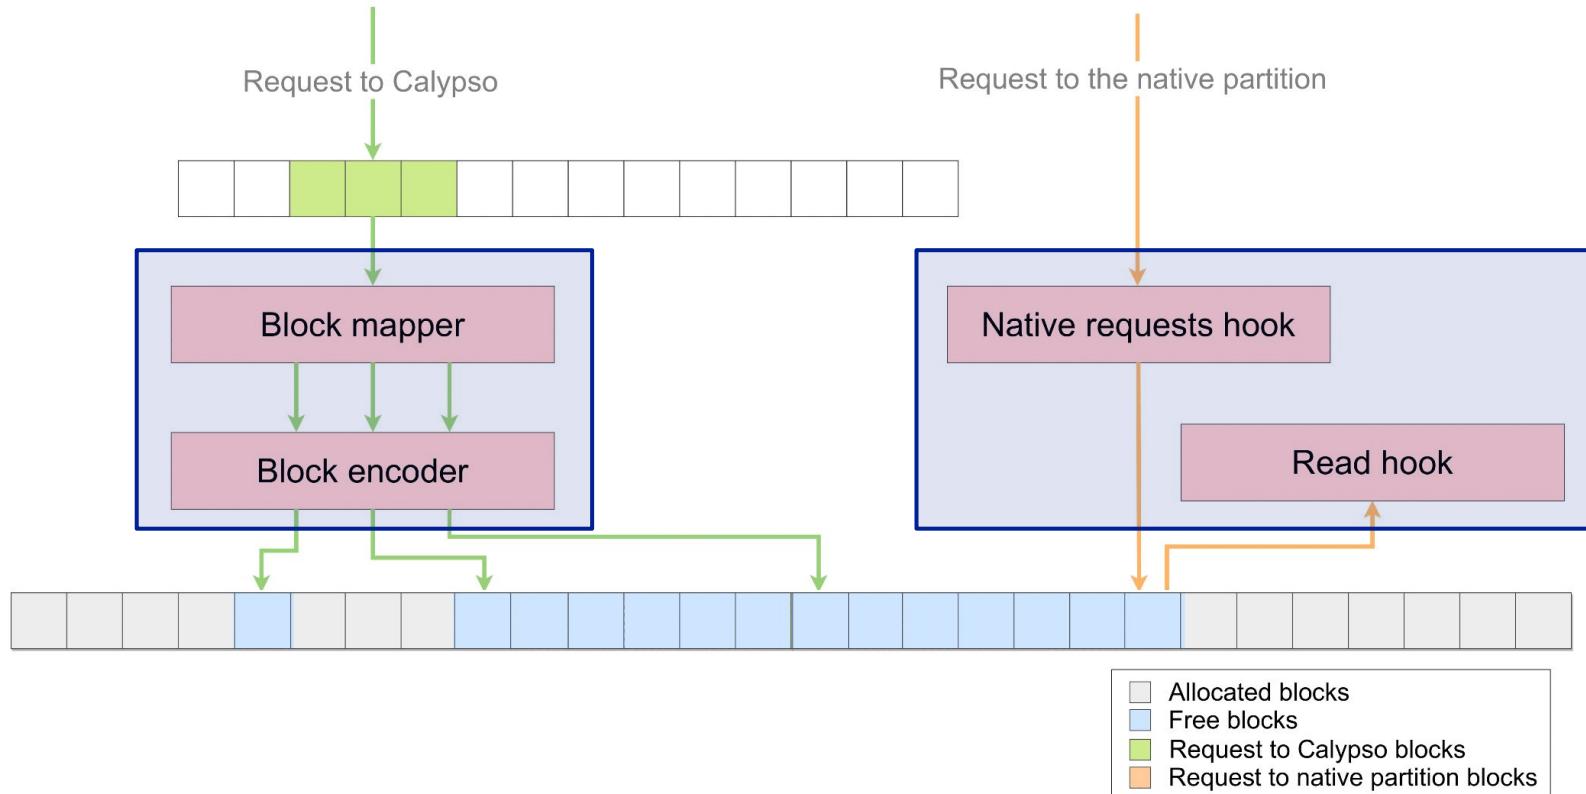


Pseudo-random
blocks change
the bit entropy
of the disk

Dependency on
external
hardware
components

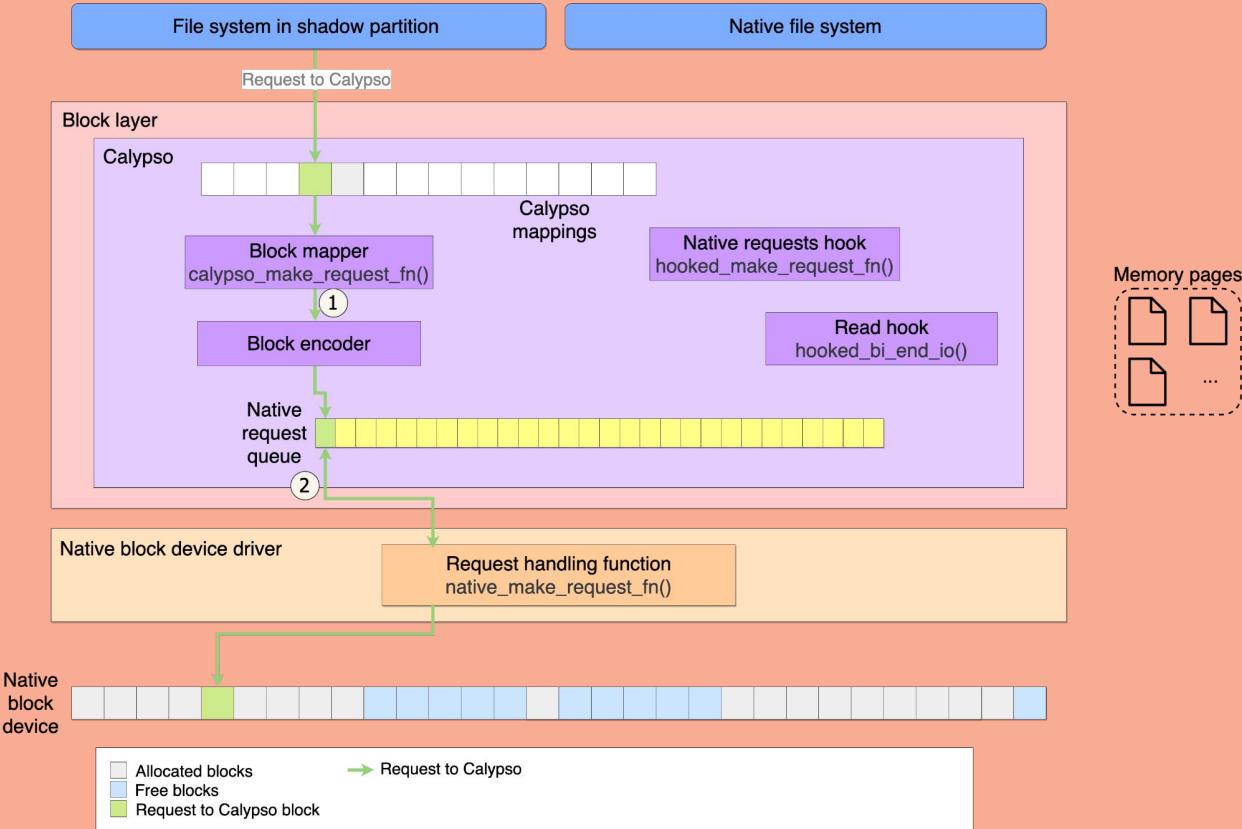


Calypso: driver



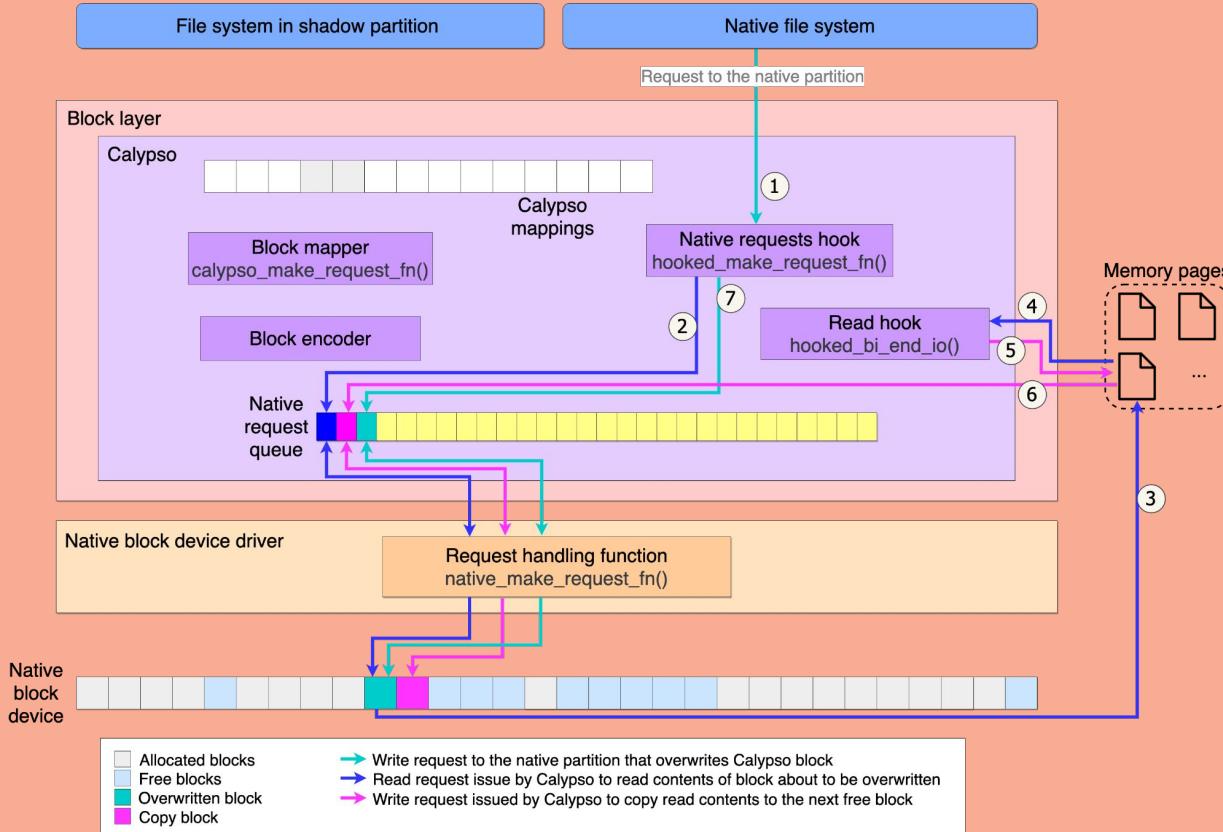


Redirecting requests

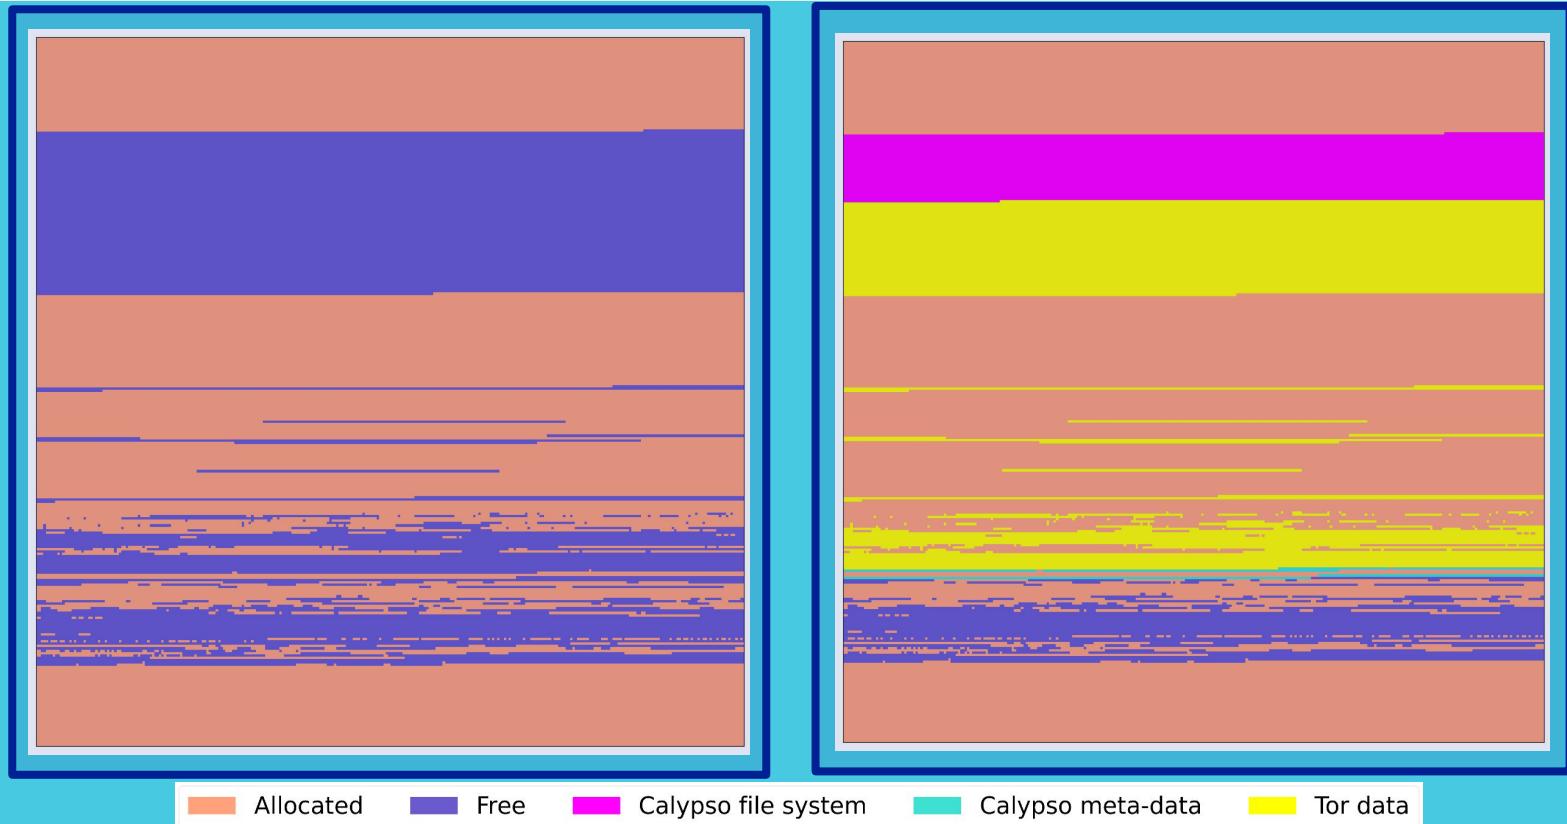




Intercepting requests and handling overwrites



Functionality



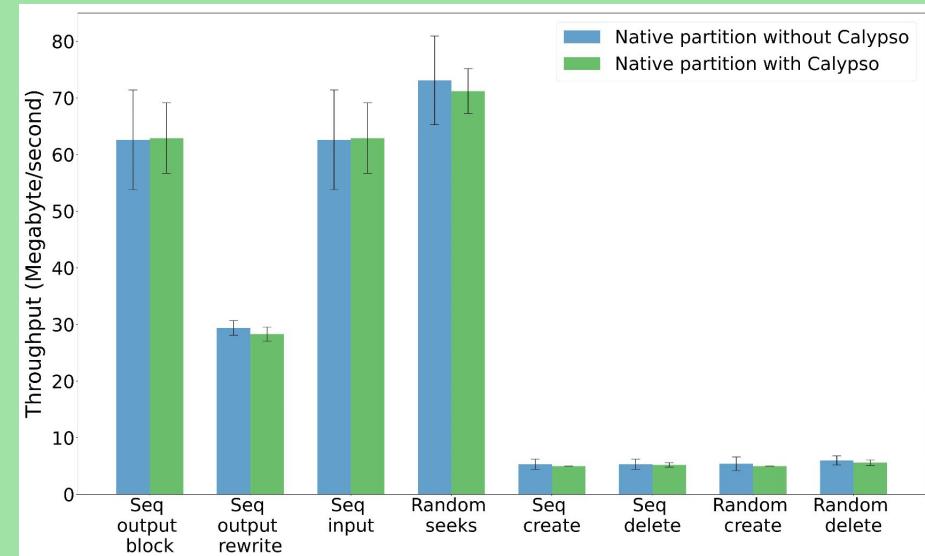
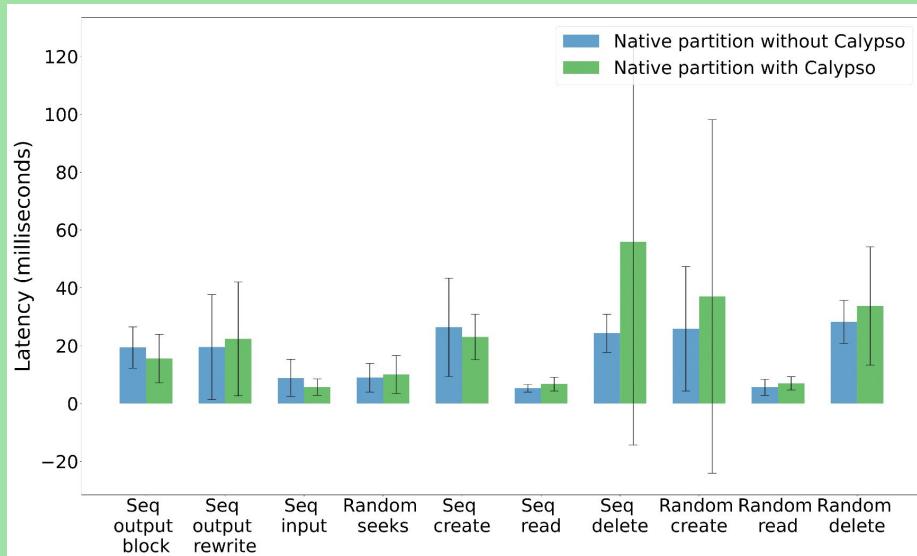


Security

Threshold	Usable blocks count	Usable blocks (%)	Storage capacity (GB)	Average differential	Minimum differential	Maximum differential
0	1048576	100.00%	4.00	0.11225	0.00000	1.00000
1	1048291	99.97%	4.00	0.11198	0.00000	0.87262
2	1048141	99.96%	4.00	0.11187	0.00000	0.74745
3	1047967	99.94%	4.00	0.11175	0.00000	0.62262
4	1047261	99.87%	3.99	0.11139	0.00000	0.49616
5	785484	74.91%	3.00	0.00306	0.00000	0.37172
6	784098	74.78%	2.99	0.00262	0.00000	0.24590
7	782983	74.67%	2.99	0.00244	0.00000	0.12039

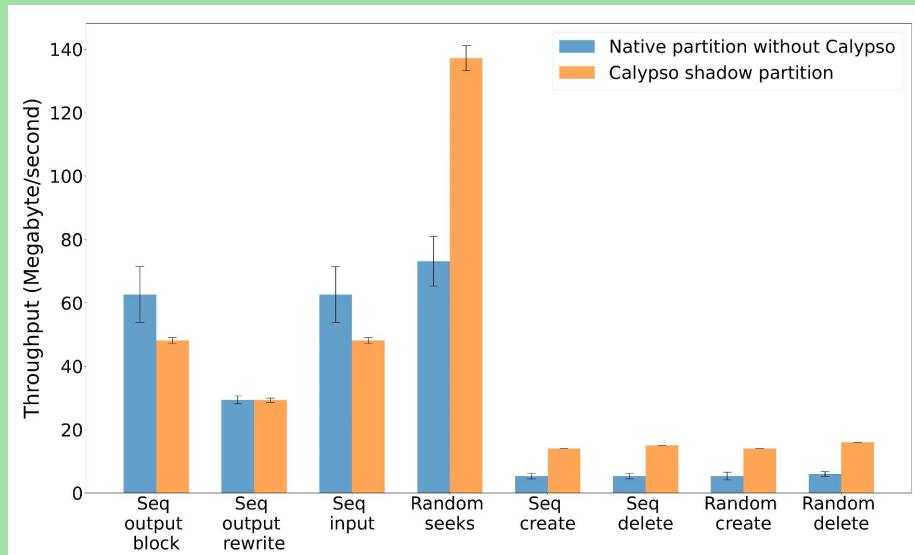
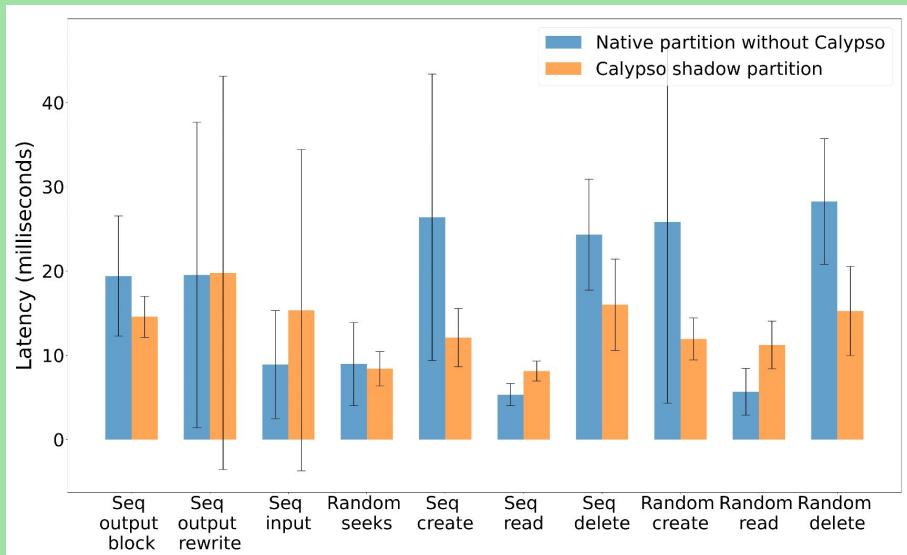


Performance



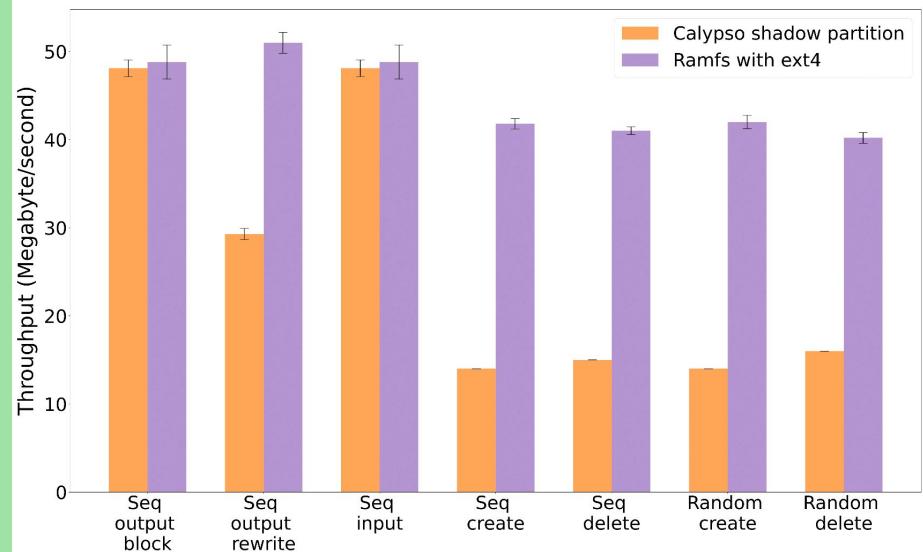
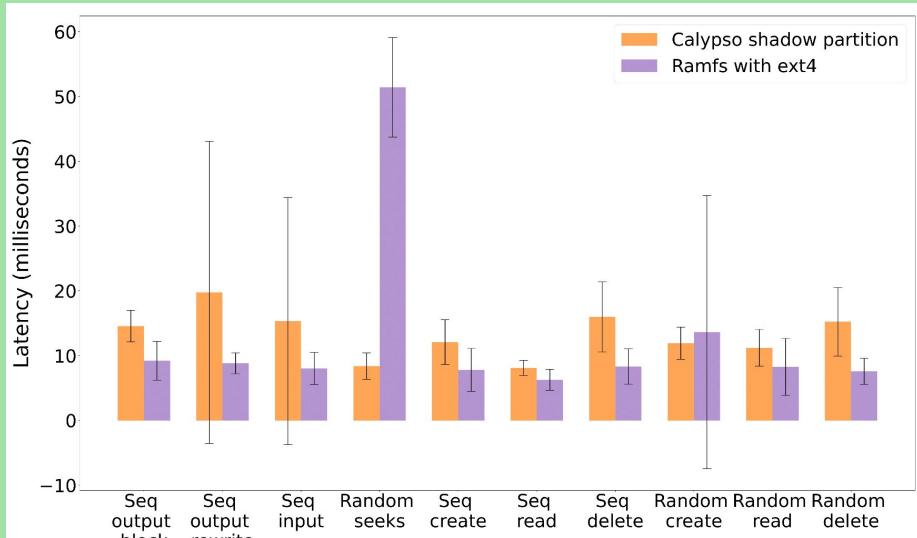


Performance





Performance





Fix slides



Contents

1. Introduction and motivation
2. **Related work**
3. Design
4. Evaluation
5. Conclusions

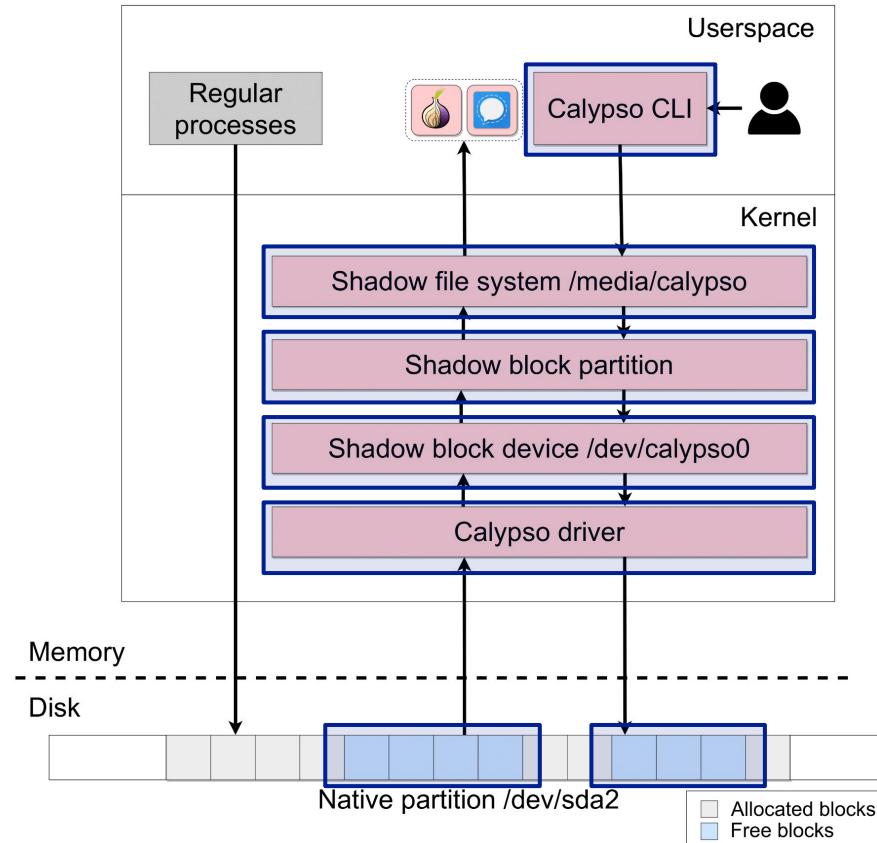


Contents

1. Introduction and motivation
2. Related work
3. Design
4. Evaluation
5. Conclusions



Calypso: architecture





Contents

1. Introduction and motivation
2. Related work
3. Design
- 4. Evaluation**
5. Conclusions



Calypso: design obstacles

Traceless bootstrap

- Encode all actions required to make Calypso available to the user
- Prevent or erases traces from executing Calypso
- Single point of failure that can compromise the system's deniability

Block allocation

- Find usable blocks without disrupting the native system or causing corruption to Calypso

Monitoring native changes to blocks

- Prevent overwrites to Calypso data
- Keep track of the native allocation of free blocks that stop being usable

Deniably encoding data blocks

- Ensure Calypso only makes deniable changes to native blocks
- Encrypted data has particularities such as high entropy. Increasing the disk entropy discloses the existence of hidden data



Calypso: design solutions

Traceless bootstrap

- Hide bootstrap files in unpredictable locations in the disk, only recoverable by the rightful user
- Erase all persistent traces of bootstrap using rootkit techniques

Block allocation

- Use the free blocks of the native file system
- Map Calypso allocated blocks with in-memory data structures:
- Keep the free or allocated status for every native block

Monitoring native changes to blocks

- Intercept all write requests to the native system to check and update the allocation in-memory data structures accordingly

Deniably encoding data blocks

- Choose only blocks with initial entropy above an established threshold



Contents

1. Introduction and motivation
2. Related work
3. Design
- 4. Implementation**
5. Evaluation
6. Conclusions



Calypso: implementation obstacles

Retrieving usable blocks for storage

- Technically challenging and file system dependant because the Linux Virtual File System does not provide an interface to retrieve the free blocks

Adapting read and write requests functionality

- **Redirecting:** passing adapted requests to the native blocks
- **Intercepting:** monitoring and preventing changes to blocks mapped by Calypso, as well as keeping track of free blocks that become allocated

Cryptography and data hiding

- Prevent adversaries from correctly retrieving Calypso's data or meta-data
- **Hiding data blocks:** generating an encryption key deterministically only with the right user's secret
- **Storing and retrieving meta-data:** finding its location without in-memory structures



Calypso: implementation solutions

Retrieve usable blocks for storage

- Iterate through the Ext4 native file system meta-data structures to retrieve all the block bitmaps into a single one.

Adapting read and write requests functionality

- **Redirecting:** adapting the received requests to map the corresponding native blocks and passing them directly to the native request queue
- **Intercepting:** hooking the native block request handling function

Cryptography and data hiding

- **Hiding data blocks:** HMAC-based extract and expand key derivation function to generate an encryption key depending on the user's password
- **Storing and retrieving meta-data:** passing a seed to a pseudo-random number generator to obtain the same sequence of numbers for a user password



Related work

Private program execution

Isolate and conceal all data generated by a program, without providing the user with plausible deniability

Deniable file systems

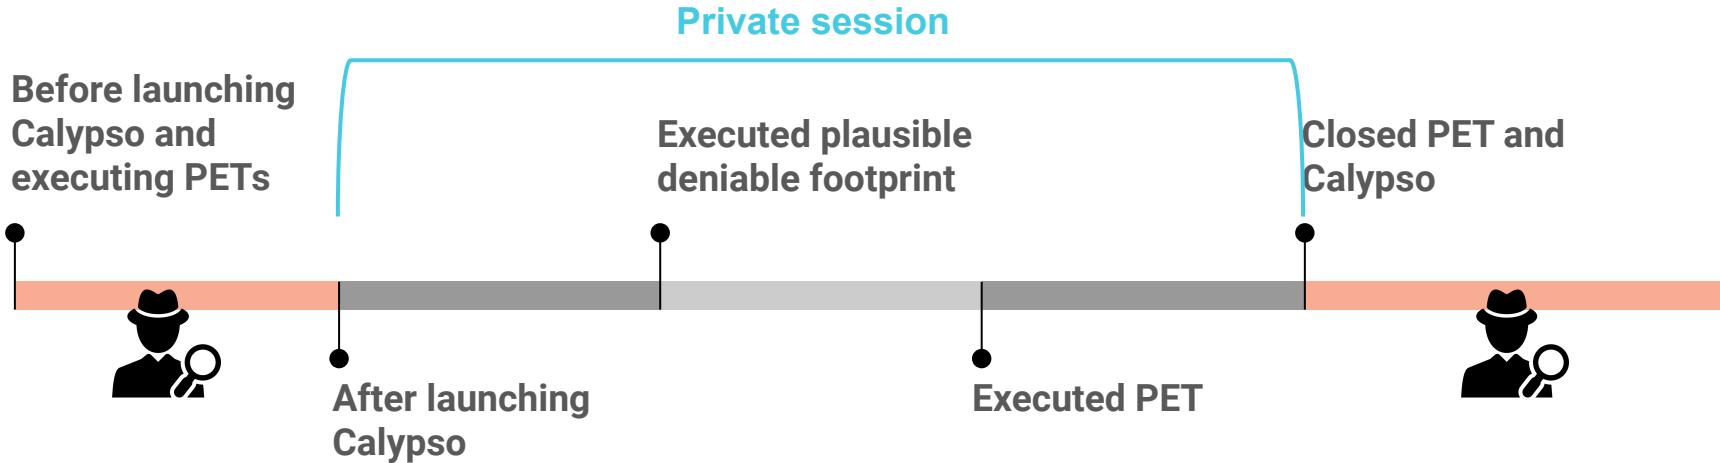
Deniable data hiding on persistent storage and single-snapshot attack prevention

Defending from multi-snapshot attacks

Multi-snapshot attack prevention by hiding modifications that indicate the existence of hidden data



Threat model



Isolation



Non-observability



Plausible deniability



Our goal: plausible deniable PET sessions

Plausible deniable PET session

John passes border control to enter the repressive country and has his device inspected



John initiates a plausible deniable PET session

John makes publications about his sensitive findings via the Tor browser on WikiLeaks

John passes border control to leave the repressive country and has his device inspected a second time

The investigator compares **Snapshot 1** and **Snapshot 2** and observes that John was the author of the publications



Plausible deniability: ability to deny knowledge or responsibility of having installed and executed PETs on their device



Conclusions

- State-of-the-art PETs are vulnerable to forensic analysis
- Calypso makes non-observable changes to disk
- Give protection against forensic analysis



Why people need to use PETs

Challenges that Internet users face:



Censorship



Lack of privacy



Surveillance

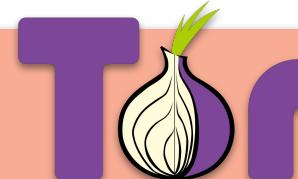
Privacy Enhancing Technologies (PETs):



Secure Messaging



VPNs



Anonymizing networks



Anti tracking



Traceless Execution Support for Privacy Enhancing Technologies

Student: Daniela Lopes

Advisor: Nuno Santos

Chairperson:

Member of the committee:

Instituto Superior Técnico



Many people are afraid of using PETs

Forensic inspection of devices



Current PETs are not designed
to remain unobservable during
forensic investigations

- Journalists
- Whistle blowers
- Travelers crossing border control

**Many people fear
the consequences
of using PETs**



Using PETs and their implications



Censorship



Lack of privacy



Surveillance

Lack of usability



Greater usage local resources



Meta-data and cryptographic artifacts

- Journalists
- Whistle blowers
- Border controls

Forensic inspection of devices

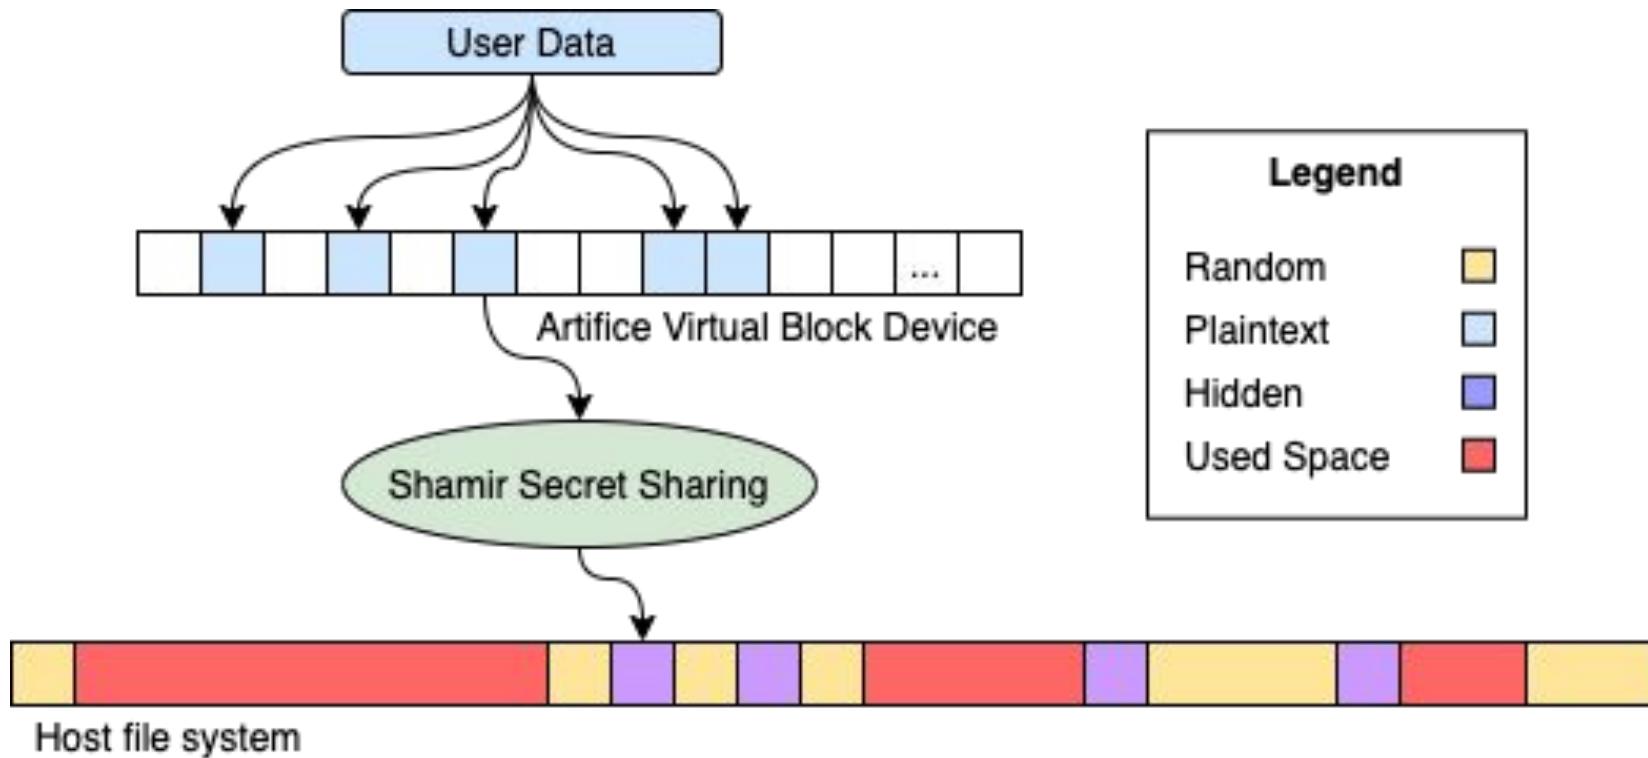


PETs do not provide plausible deniability

Many people fear the consequences of using PETs



Artifice: encoding user data





Artifice: overview

Protect users in life-threatening situations



Provide a deniable way to store sensitive data

Adversary is unable to distinguish blocks encoding hidden data from blocks with random data



Artifice is not deniable

— Pseudo-random blocks change the bit entropy of the disk

— Dependency on external hardware components

Artifice is not deniable!



Proposed solution: Calypso

Shadow drive

- Secondary file system that uses the unused blocks of the native file system



Shadow Containers

- Isolated environment to execute programs deniably

No observable changes are made to the persistent state



Calypso: design obstacles

Traceless bootstrap

Block allocation

- Finding and maximizing usable space
- Encoding data within the blocks while maintaining the entropy of the disk
- Fault tolerance and data loss

Monitoring native changes to blocks

-



Conclusions

State-of-the-art PETs are
vulnerable to forensic
analysis

Calypso offers isolation
and maintains the entropy
of the disk

Provide a traceless
execution environment to
increase PET usage by
censored users

Give protection against
forensic analysis



Using PETs and their implications



Censorship



Lack of privacy



Surveillance

Forensic inspection
of devices



PETs do not provide
plausible deniability

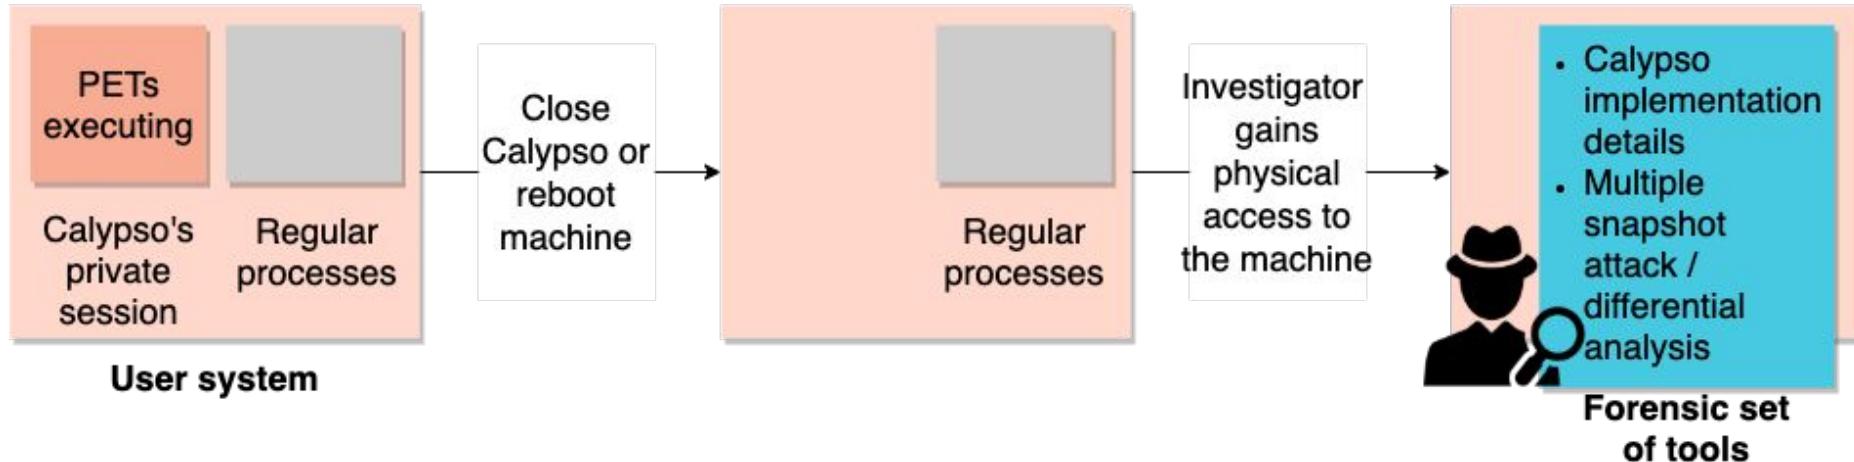


Many people fear
the consequences
of using PETs!

- Journalists
- Whistle blowers
- Travelers passing border controls



Threat model





Using PETs and their implications



Censorship



Lack of privacy



Surveillance

Forensic inspection
of devices



PETs do not provide
plausible deniability



Many people fear
the consequences
of using PETs!

- Journalists
- Whistle blowers
- Travelers passing border controls



Goals

Challenges

Environment to execute PETs
deniably with persistent state

Maintain a PET's functionality
and performance

Solutions

Parasite on the free blocks of the
native system

Preserve the initial entropy
characteristics of the disk



Retrieving usable blocks for storage



Cryptography and data hiding



Contents

1. Introduction and motivation
2. Related work
3. Design
4. Implementation
5. Evaluation
6. Conclusions



Why people are afraid of using PETs

- Lack of usability
- Require large amounts of local resources
- Leaves an extensive footprint



Forensic inspection
of devices



Why people are afraid of using PETs

Forensic inspection
of devices



- Journalists
- Whistle blowers
- Travelers passing border controls

PETs do not provide
plausible deniability

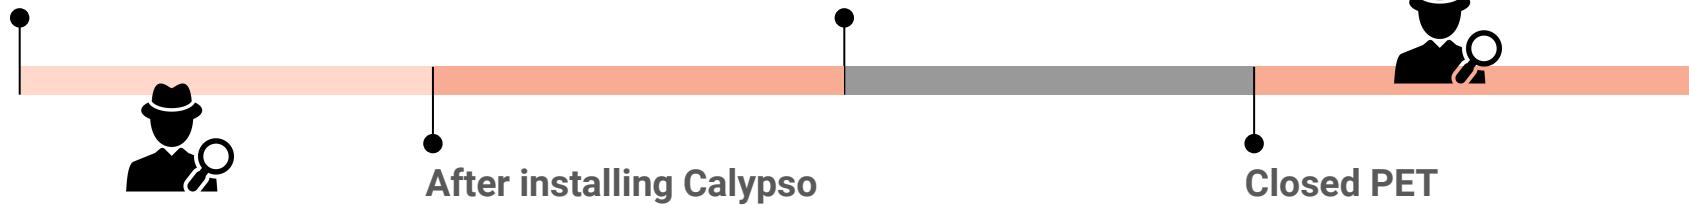


Many people fear
the consequences
of using PETs!



Threat model

Before installing Calypso
and executing PETs



After installing Calypso

Executed PET

Closed PET

Isolation



Non-observability



Plausible deniability



Traceless bootstrap



Block allocation



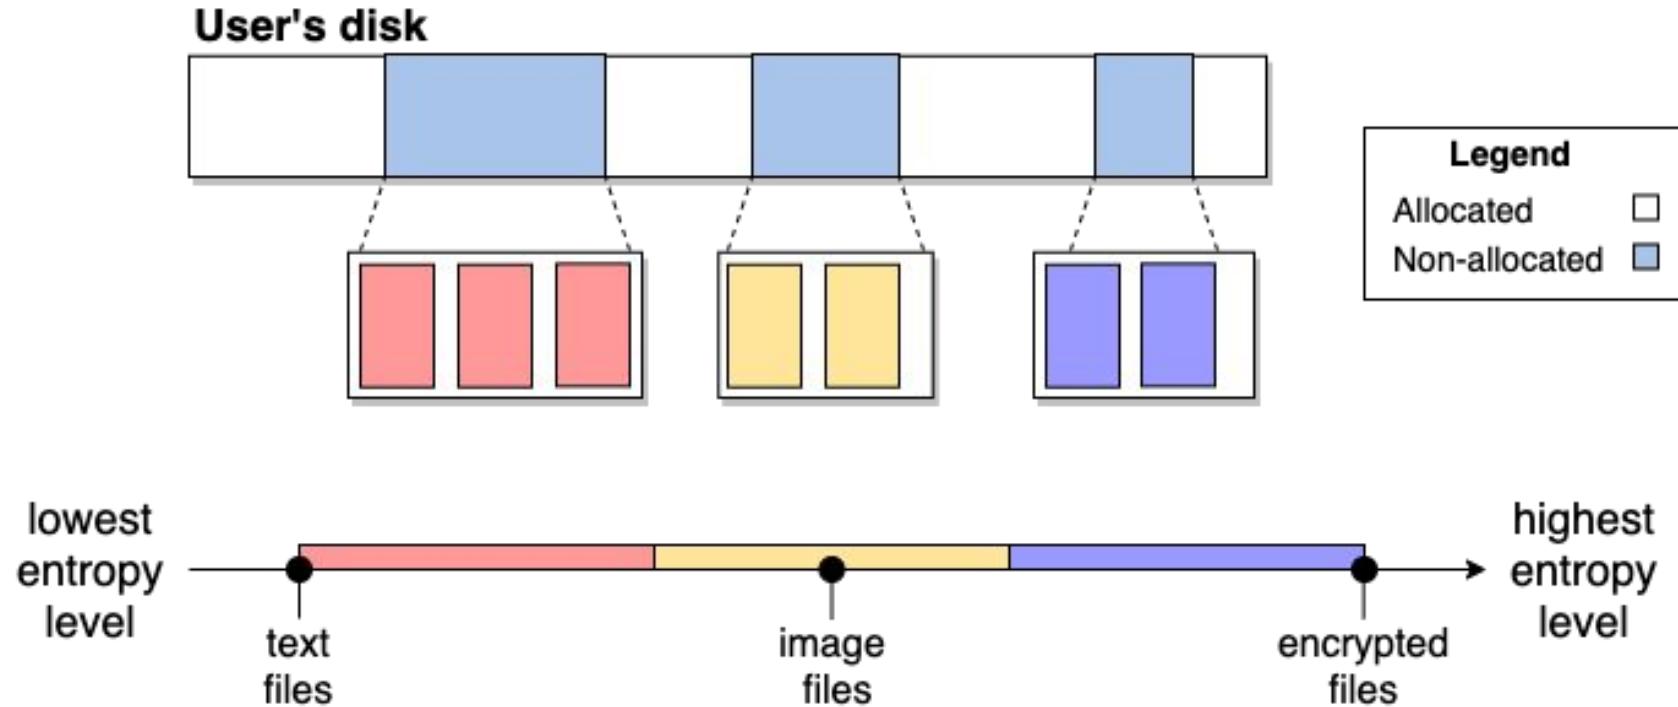
Monitoring native changes to blocks



Deniably encoding data blocks



Profiling the user's disk





Isolation

File category	Initial state		Tor native		Tor Calypso	
	Accesses count	Different file count	Accesses count	Different file count	Accesses count	Different file count
Temporary memory only	2138	67	4268	742	11377	5298
Temporary in disk	124	2	99	54	618	26
System config and OS resources	448	27	3015	161	2913	145
Binaries, libs and apps	239	62	3276	1389	1516	370
User data	0	0	2553	510	109	63
Calypso	0	0	0	0	318	75
Other	200	1	240	37	421	27



Evaluation

Functionality

Isolation

Security

Performance



Calypso: design obstacles

Functionality

- Test Calypso's ability to sustain the execution of PET applications without file system errors

Isolation

- Ensure Calypso leaves no persistent traces

Security

-

Performance

-



Why people are afraid of using PETs

Lack of usability

Large amounts of local resources

Extensive digital footprint

Forensic inspection of devices

PETs do not provide plausible deniability

People are afraid of using PETs



Goals and challenges

Goal

Build a system to execute PETs with plausible deniability, allowing users to perform private sessions

Challenges

Make all changes to persistent state non-observable

Maintain PETs usability, functionality and performance

Solutions

Shadow partition that uses the free blocks of the native file system

Encoded data can only be recovered by the rightful user



Contributions

Design

Calypso, a steganographic storage system

- Parasites on the free blocks of the native system to conceal data
- Preserves the initial entropy characteristics of the disk

Implementation

Fully functional prototype as a Linux kernel module for version 5.4

Evaluation

Extensive evaluation of the prototype



Conclusions

State-of-the-art PETs are vulnerable to forensic analysis

Calypso makes non-observable changes to disk

Give protection against forensic analysis